


Announcements



- **Weekly Reading Assignments:
Chapters 23 & 24 (CLRS)**

Definitions for Shortest Paths



- Think of vertices as cities and the edge weights as the distance from one city to another. Define the *length* of a path to be the sum of edge weights along the path. Define the *distance* between two vertices, u and v , $\delta(u,v)$ to be the length of the minimum length path from u to v .
- A *shortest path* from vertex u to vertex v is then defined as any path p with weight $w(p) = \delta(u,v)$.

Single-Source Shortest Paths



- Given a directed graph $G = (V, E)$ with edge weights and a distinguished **source vertex**, $s \in V$, determine the distance from the source vertex to every vertex in the graph.
- BFS finds short-paths from a single source vertex to all other vertices in $O(n+e)$ time, assuming the graph has no edge weights.
- Edge weights can be negative; but in order for the problem to be well-defined there must be no cycle whose total cost is negative.

Variants



- **Single-destination shortest-paths problem:** Find a shortest path to a given destination vertex t from every vertex $v \in V$.
- **Single-pair shortest-path problem:** Find a shortest path from u to v for given vertices u and v . If we solve the single-source shortest paths problem, we also solve this problem.
- **All-pairs shortest-paths problem:** Find a shortest path from u to v for every pair of vertices u and v . This can be solved by the single-source problem run for each vertex.

Relaxation



- Maintain an estimate of the shortest path for each vertex v , call it $d[v]$.
- Initially $d[v]$ will be the length of the shortest path that the algorithm knows from s to v . This value will always be greater than or equal to the true shortest path distance from s to v .
- Initially, we know of no paths, so all $d[v]=\infty$ & $d[s]=0$.
- As the algorithm goes on and sees more vertices, it tries to update $d[v]$ for each vertex in the graph, until all $d[v]$ values converge to true shortest distances.

Find Shortest Path by Relaxation



- If the solution is not yet an optimal value, then push a little closer to the optimum. If we find a path from s to v shorter than $d[v]$, then update $d[v]$.
- Consider an edge from a vertex u to v whose weight is $w(u,v)$. Suppose that we have already computed current estimates on $d[u]$ and $d[v]$. We know that there is a path from s to u of weight $d[u]$. By taking this path and following it with the edge (u,v) we get a path to v of length $d[u] + w(u,v)$. If this path is better than the existing path of length $d[v]$ to v , we should take it.

Relax (u, v)



1. if $d[u] + w(u, v) < d[v]$ // is the path thru u shorter?
2. then $d[v] \leftarrow d[u] + w(u, v);$ // yes, then take it.
3. $\pi[v] \leftarrow u;$
 // the shortest way back to the source is thru
 // u by updating the predecessor pointer

NOTE: If we perform Relax (u, v) repeatedly over all edges of the graph, all the $d[v]$ values will eventually converge to the true final distance values from s . How to do this most efficiently?

Dijkstra's Algorithm




- Maintains a subset of vertices, $S \subseteq V$, for which we know their true distance $d[u] = \delta(s, u)$. Initially $S = \emptyset$ and we set $d[s] = 0$ and all others to ∞ . One by one we select vertices from $V - S$ to add to S .
- For each vertex $u \in V - S$, we have computed a distance estimate $d[u]$. The greedy approach is to take the vertex for which $d[u]$ is minimum, i.e. take unprocessed vertex that is closest to s .
- We store the vertices of $V - S$ in a priority queue (heap), where the key value of each vertex u is $d[u]$. All operations can be done in $O(\lg n)$ time.


Dijkstra(G, w, s)




1. $Q \leftarrow V[G]$ and $S \leftarrow \emptyset$
2. for each vertex $u \in Q$ // initialization: $O(V)$ time
3. do $d[u] \leftarrow \infty$ and $\pi[u] \leftarrow \text{NIL}$
4. $d[s] \leftarrow 0$ // start at the source
5. $\pi[s] \leftarrow \text{NIL}$ // set parent of s to be NIL
6. while $Q \neq \emptyset$ // till all vertices processed
7. do $u \leftarrow \text{Extract-Min}(Q)$ // select closest to s
 $S \leftarrow S \cup \{u\}$
8. for each $v \in \text{adj}[u]$
9. do if $v \in Q$ and $(d[u] + w(u,v) < d[v])$
10. then $\pi[v] \leftarrow u$
11. $d[v] \leftarrow d[u] + w(u,v)$ // Relax (u,v)
12. decrease_Key($Q, v, d[v]$)




Example: Dijkstra's Algorithm




Black: in S




Example: Dijkstra's Algorithm




Black: in S




Example: Dijkstra's Algorithm




Black: in S




Example: Dijkstra's Algorithm




Black: in S




Example: Dijkstra's Algorithm



Black: in S




Example: Dijkstra's Algorithm




Black: in S


Example: Dijkstra's Algorithm




Black: in S




Example: Dijkstra's Algorithm



Black: in S

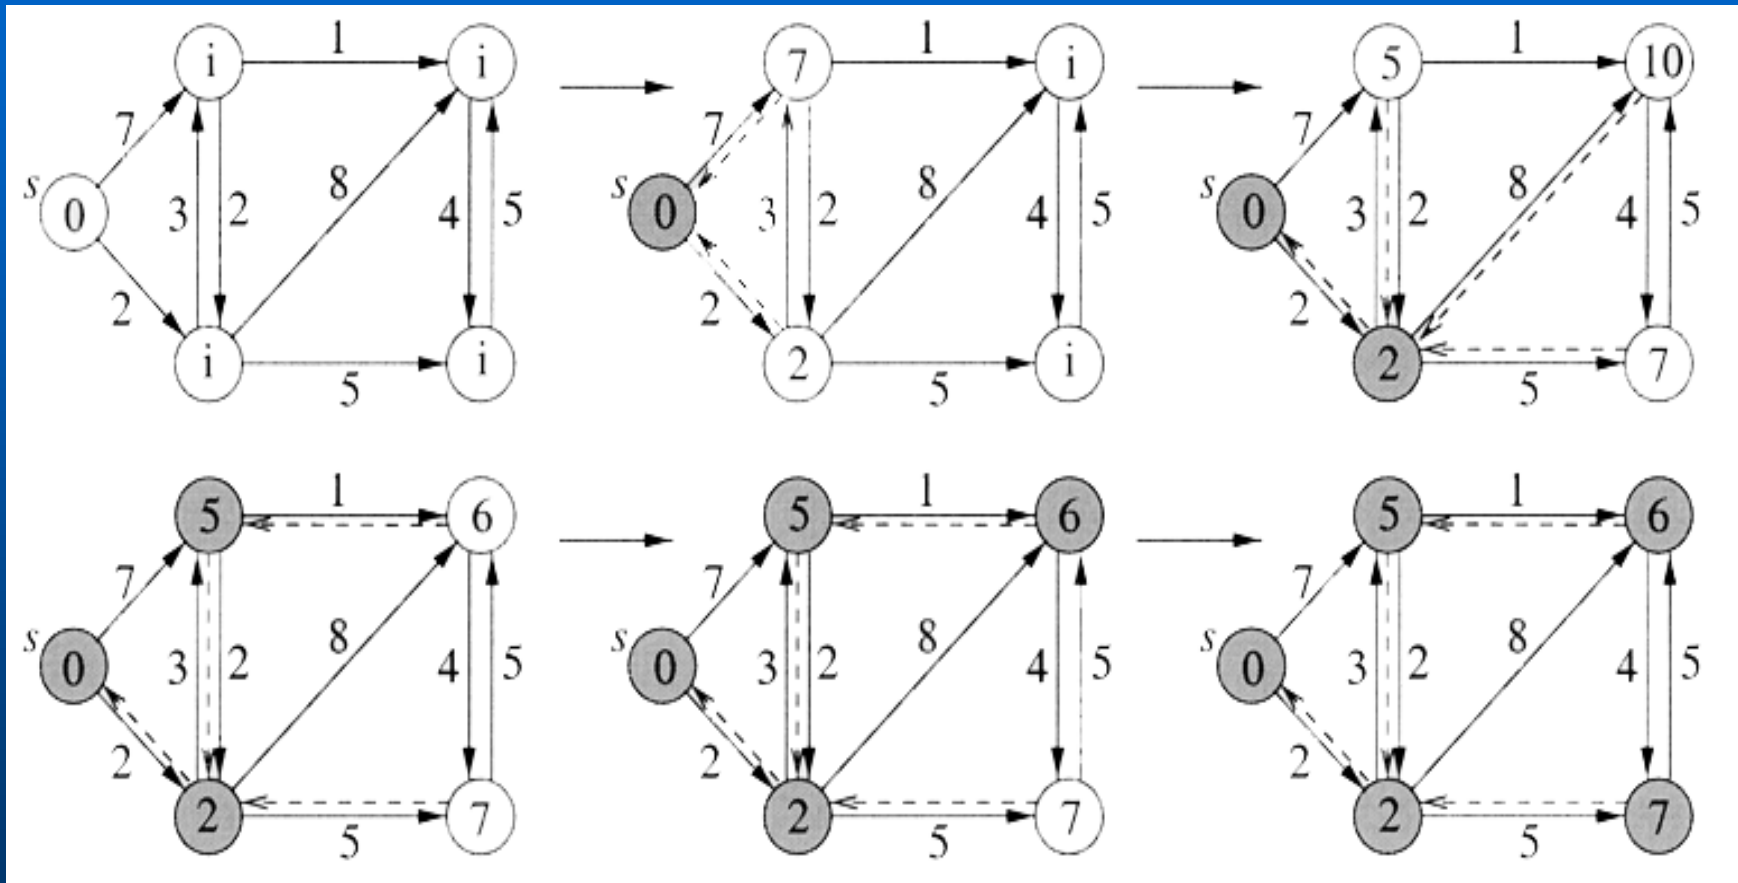


Example: Dijkstra's Algorithm



Black: in S

The Sequence of Relaxations



Correctness



- Let $\delta(s,v)$ be length of true shortest path from s to v .
Need to show: $d[v]=\delta(s,v)$ for each v when the algorithm terminates
- Lemma: When a vertex u is added to S , $d[u]=\delta(s,u)$.


Proof:

- Suppose that at some point Dijkstra's algorithm first attempts to add a vertex u to S for which $d[u] \neq \delta(s,u)$.
 - Note $d[u]$ is never less than $\delta(s,u)$, thus $d[u] > \delta(s,u)$.

Proof (I)



- Just prior to the insertion of u , consider the true shortest path from s to u .
 - Because $s \in S$ and $u \in V - S$, at some point this path must first jump out of S .
 - Let (x, y) be the edge where it jumps out, so that $x \in S$ and $y \in V - S$. (It might happen that $x=s$ and/or $y=u$).



Proof (II): $y \neq u$



- We argue that $y \neq u$ after all
 - Since $x \in S$ we have $d[x] = \delta(s, x)$. (Remember that u was the first vertex added to S that violated this criterion.)
 - Since we applied relaxation to x when it was added, we would have set $d[y] = d[x] + w(x, y) = \delta(s, y)$.
 - Since (x, y) is on a shortest path from s to u , it is on a shortest path from s to y . Thus $d[y]$ is now correct.
 - By hypothesis, $d[u]$ is not correct, so u and y cannot be the same.

Conclusion of Proof



- Notice:
 - y appears somewhere along the shortest path from s to u (but not at u)
 - All subsequent edges following y are of weight ≥ 0 ,
- Therefore, $\delta(s,y) \leq \delta(s,u)$, so $d[y] = \delta(s,y) \leq \delta(s,u) < d[u]$
- But, $d[u] \leq d[y]$ because u is added before y , and we add nodes with lower d values first.

Contradiction!

Correctness



- We have just proved that, when a vertex u is added to S , $d[u] = \delta(s, u)$.
- Every vertex is eventually added to S , so the algorithm assigns the correct distances to all vertices.
- This guarantees that the shortest-paths tree is correct because relaxation correctly updates parent pointers (see book for details).