


Longest Prefix Matching: Applying Algorithms to Networking

Ketan Mayer-Patel

October 11, 2005


What's a network?

- A collection of computers connected together.



What's an internetwork?

- A collection of networks connected together.



Routing and Forwarding

- So, how does a message get from computer A, to computer B?
 - Think of a letter.
 - What do you need to get a letter from point A to point B?
 - Address
 - What makes up an address?
 - Street and number
 - City, State, and ZipCode.

Routing and Forwarding

- We can think of the letter being delivered in two parts:
 - First, we get it to the right zipcode.
 - Then, we get it to the right house.
- This happens for computer networks as well.
 - First we get it to the right network.
 - Then we get it to the right machine.

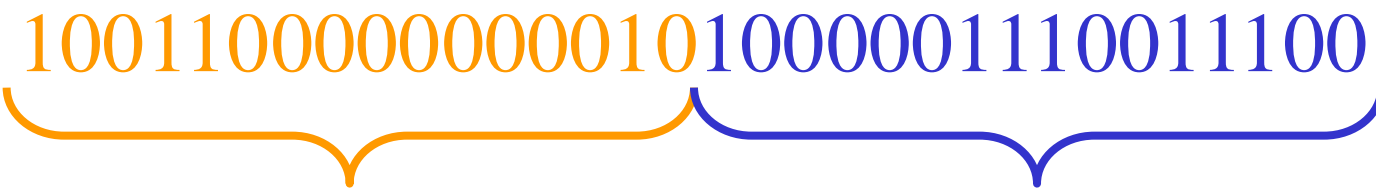
IP Addresses

- A computer's address on the Internet:
 - 32 bits long
 - Typically written in “dot notation”
 - 152.2.131.156
 - It has two parts:
 - Network address
 - Host address
 - Number of bits used for each part is variable.

IP Addresses

- To understand an IP address completely, need to know:
 - 32-bit address
 - How many bits are for the network.

- Example: 152.2.131.156

- 
100110000000000101000001110011100
Network Number Host Number

Forwarding

- In the middle of the Internet, only the network part of an address matters.
- Message comes in some incoming link.
- Need to decide on which outgoing link to send the message on.
- This is called hop-by-hop forwarding.

Routing Table

- To do forwarding:
 - Build a table for all network numbers.
 - Table associated net number with outgoing link.
 - When message arrives, look up network number, find associated link, send message.
- Building the routing table is a different problem.

The Problem with Scale

- What do we care about?
 - Speed.
 - Want to make forwarding look up as fast as possible.
 - Space
 - The number of network number is quite large.
 - If we have to have a line in the table for every possible network number, then each router would need a large amount of memory.
 - Also affects speed.

Longest Prefix Routing

- Insight: we can compress a lookup table by dealing with ranges instead of individual network numbers.

Longest Prefix Example

000	A	00*	A, 2	0*	A, 1	*	A, 0
001	A	010	A, 3	011	B, 3	011	B, 3
		011	B, 3	1*	B, 1	1*	B, 1
010	A	100	A, 3	100	A, 3	100	A, 3
011	B	101	B, 3				
		11*	B, 2				
100	A						
101	B						
110	B						
111	B						

Problem Recap

- Given:
 - A set of prefixes.
 - Length of each prefix.
 - An address to match.
- Want:
 - Longest matching prefix.

Linear Brute Search

- Examine each entry.
- Remember the longest match.
- Time for adding an entry?
 - $O(1)$
- Time for deleting an entry?
 - $O(n)$
- Avg. time for lookup?
 - $O(n/2) \sim O(n)$
- Worst case lookup?
 - $O(n)$
- Space needed?
 - $O(n)$


Sorted Ranges

- From each table entry, generate two “markers”.
 - Each marker is one bit longer than longest.
 - Left marker created by filling in the rest with 0.
 - Right marker created by filling in the rest with 1.
 - Associate prefix length and prefix with markers.
 - Sort markers (break ties with length).

Sorted Range Example

* A, 0
011 B, 3
1* B, 1
100 A, 3

Markers have 4 bits.



Sorted Range Example

- Once sorted we can use binary search to do lookup.

0000 <i>A</i> , 0	0110 <i>B</i> , 3	0111 <i>B</i> , 3	1000 <i>B</i> , 1	1000 <i>A</i> , 3	1001 <i>A</i> , 3	1111 <i>B</i> , 1	1111 <i>A</i> , 0
[[]	[[]]]

Sorted Range Performance


- Time for adding an entry?
 - $O(n \log n)$
- Time for deleting an entry?
 - $O(n \log n)$
- Time for lookup?
 - $O(\log n)$
- Space needed?
 - $O(2n)$

Radix Tree

- Popular in actual routers.
- Build a binary tree.
- Each level of the tree is indexed by next bit.
- Only build as much of the tree as needed for all of the entries.
- Associate link with each node that matches a prefix in the tree.


Radix Tree Example

*	A, 0
011	B, 3
1*	B, 1
100	A, 3



Radix Tree Example

*	A, 0
011	B, 3
1*	B, 1
100	A, 3



Radix Tree Performance

- Time for adding an entry?
 - $O(\text{entry length}) = O(\log e)$
- Time for deleting an entry?
 - $O(\log e)$
- Time for lookup?
 - $O(\log e)$
- Space needed?
 - $O(2^{\log n+1}) \Rightarrow$ Worst case, usually much better.

One more question.

- Still have the problem of building the prefix table to begin with.
- Which of these data structures and algorithms can help us with that?
 - Radix tree.